

npm registry

dev-ops deep-dive




C J Silverio

director of engineering, npm

@ceejbot

registry 1.0

embedded in couchdb



javascript

but not node

the shame, the shame

advantages

- hey! it was a simple working system
- couchdb's replication made mirrors easy
- didn't have to implement auth
- got away with storing package tarballs as couch attachments
- worked for a longer time than we deserved

disadvantages

- all of this fell over at scale
- tarballs fell over first
- we aren't erlang experts
- not modular; hard to work on

late 2013: stay up

- pulled out tarballs into Joyent Manta**
- put varnish in front of everything**
- fastly CDN for geolocality**

early 2014: stability

- tarballs onto a file system**
- found & stomped problems with our couchdb installation**
- load-balanced everything**
- operational maturity**
- big sign of success: many mirrors shut down**

now we're stable!

npm's next goal:

be self-sustaining

end 2014: rewrite

- we are node experts!**
- microservices: node's natural architecture**
- future scaling**
- ability to add features easily**
- scoped modules!**

scoped modules aka namespaces

- **hyperfs: the famous module**
- **@mikeal/hyperfs: super-hip fork**
- **@ceejbot/hyperfs: my completely unrelated private module**

Everybody can make public scoped modules. \$7/month and you can create private scoped modules.

team

- **3 engineers on the registry & operations**
- **2 engineers on the website**
- **2 engineers on the command-line client**

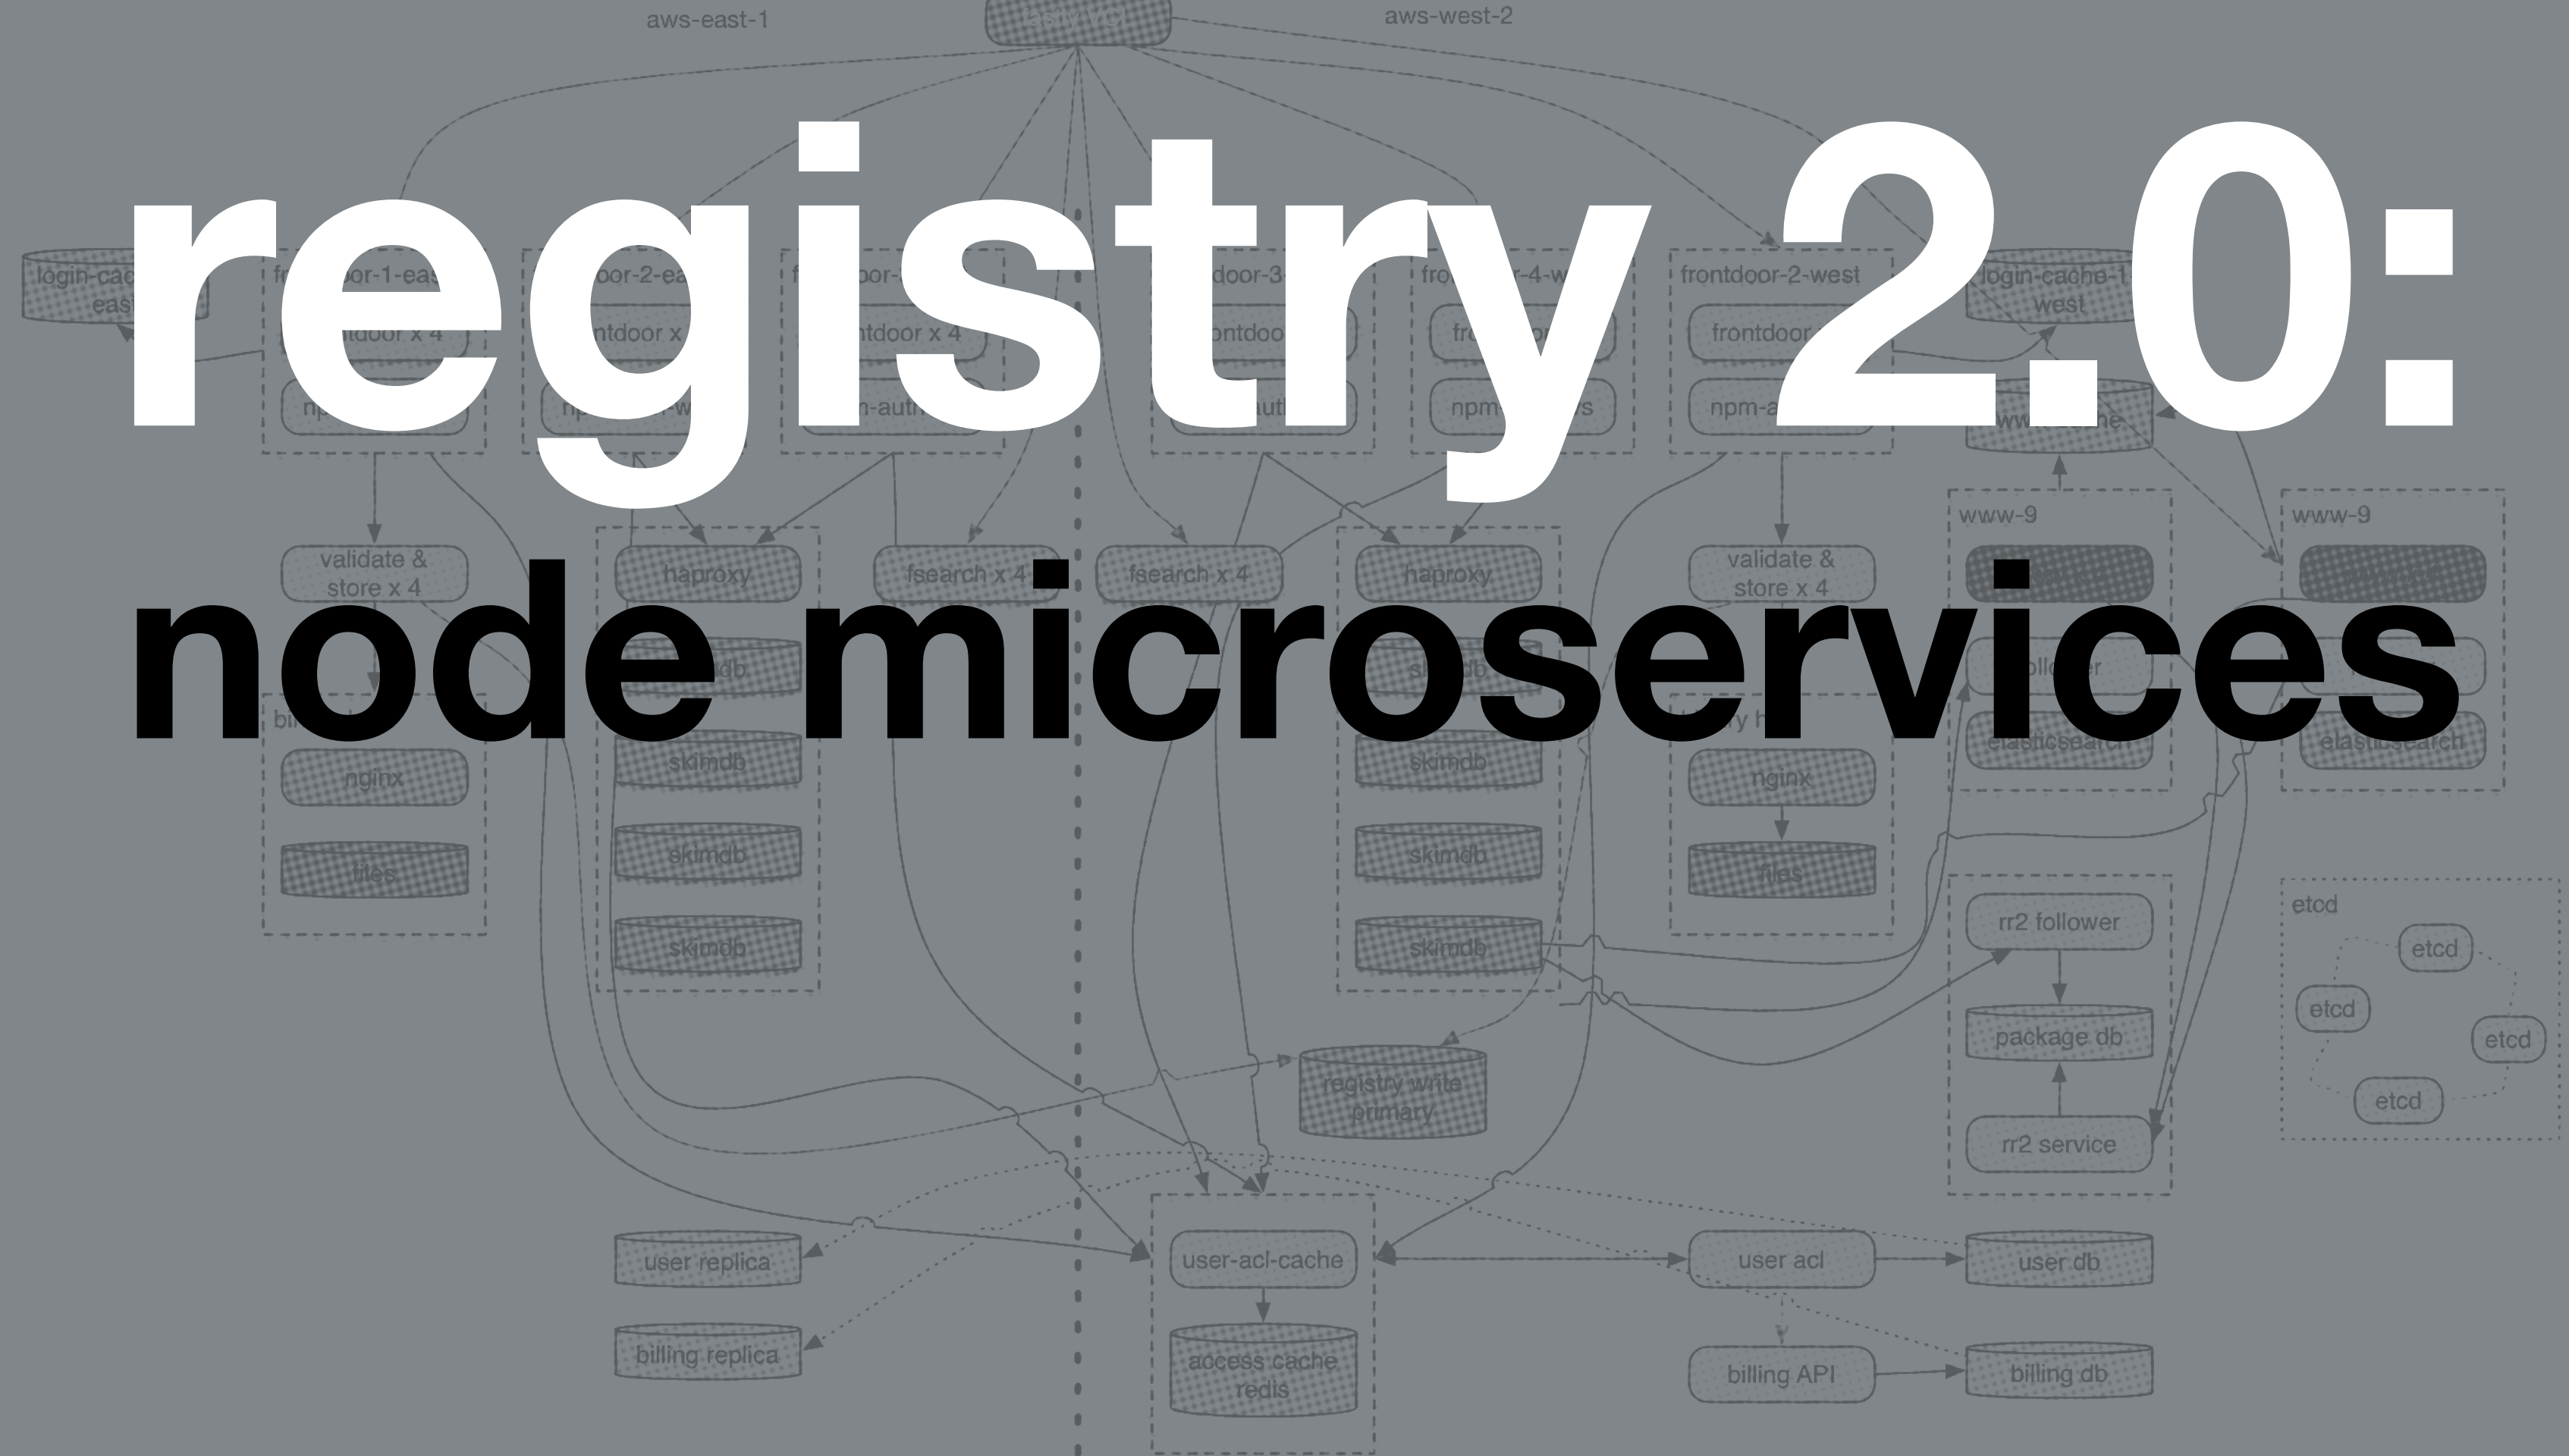
shipp^{ed} the core of it
as **npm-enterprise**
"npm in a box" service
(our other way to make \$)

had a working registry in node
before we migrated the
public registry to it

in production April 2015
scoped modules
were a feature flip

registry 2.0:

node microservices



the stack (top)

- **Fastly as our CDN (faster in Europe!)**
- **AWS EC2**
- **Ubuntu Trusty**
- **nagios + PagerDuty**
- **Github hosts our code**
- **TravisCI for public & private repos**

the stack (middle)

- haproxy for load balancing & tls termination**
- a couple instances of pound for tls (legacy)**
- nginx for static files**
- redis for caching**

the databases

- **couchdb for package data storage**
- **postgres for users, billing, access control lists**
- **replica of the package data in postgres to drive website**

big node modules

- **web site only: hapi**
- **everything else: restify**
- **knex to help with postgres**

restify

- barely a framework
- trivial to get a json api running
- observable
- sinatra/express routing
- we like the connect middleware style

conventions across services

- monitoring endpoints same for all
- every process has a repl
- json logging
- config mostly through cmd-line arguments
- some environment variable passing

configuration via etcd

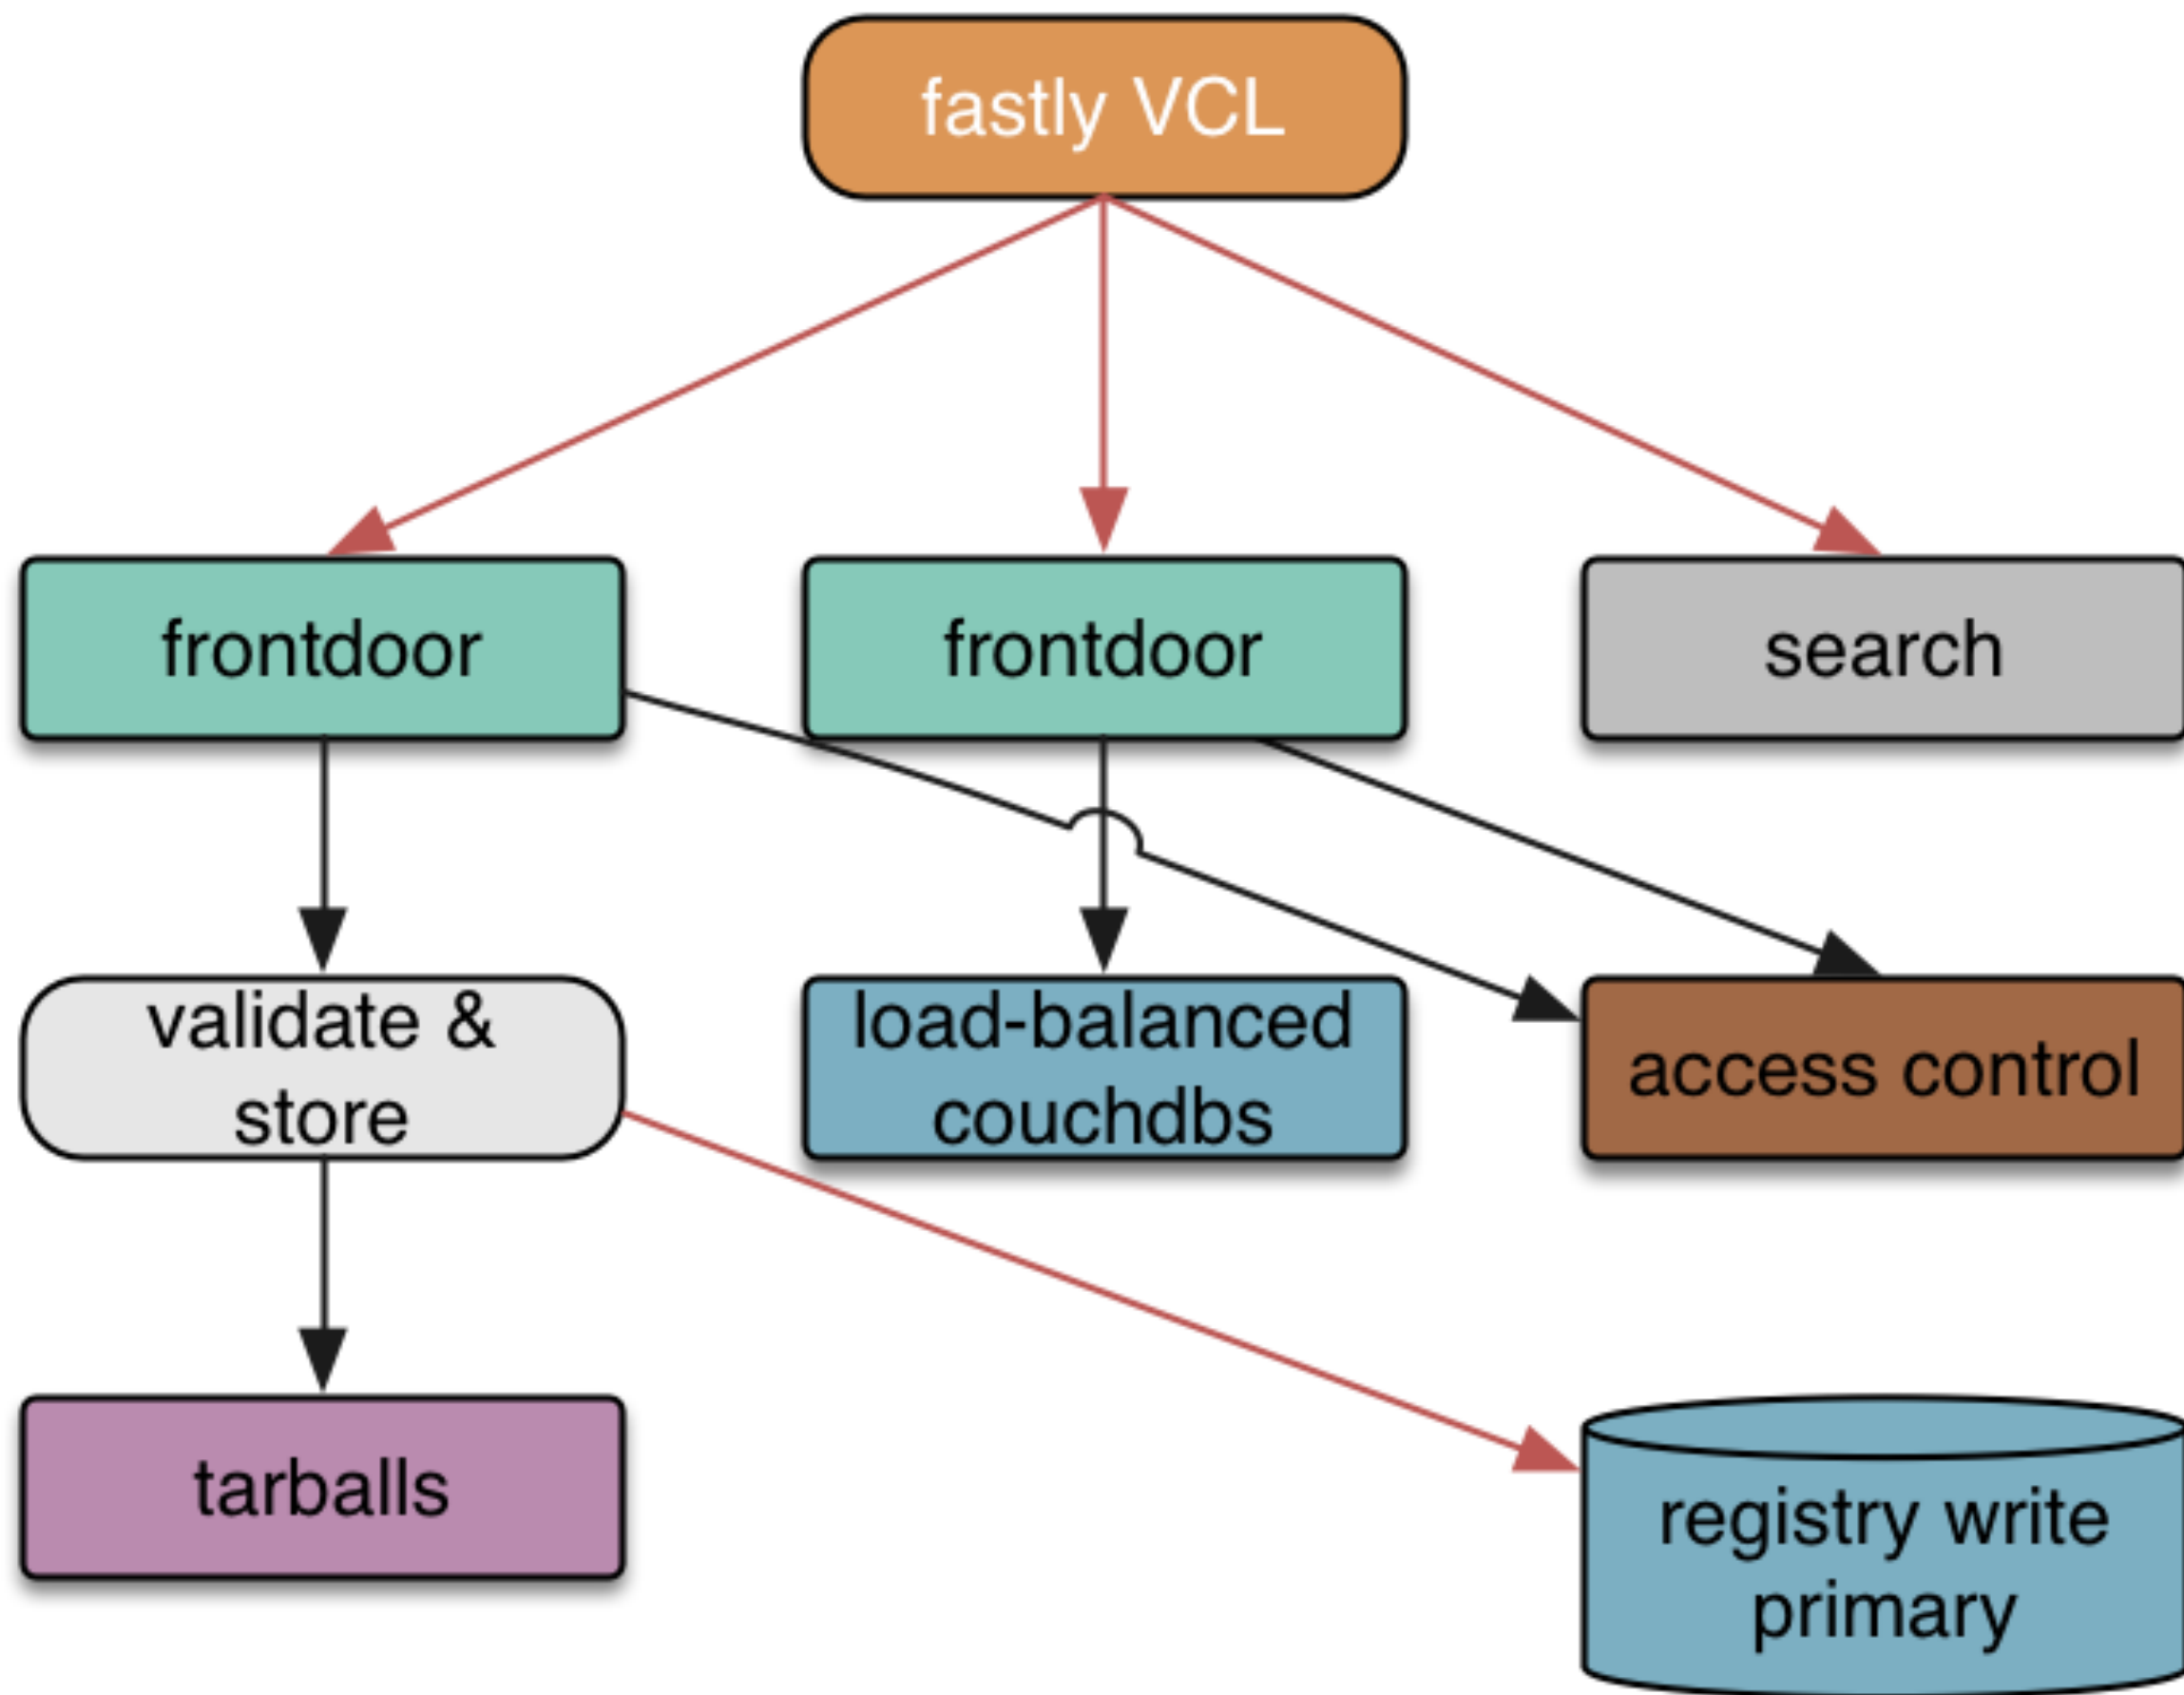
<https://github.com/coreos/etcd>

A highly available key/value store intended for config & service discovery. We recursively store & extract json blobs from it using `renv`.

`ndm` tool transforms json into command-line options in an upstart script.

automation via ansible
any box can be replaced
by running an ansible play

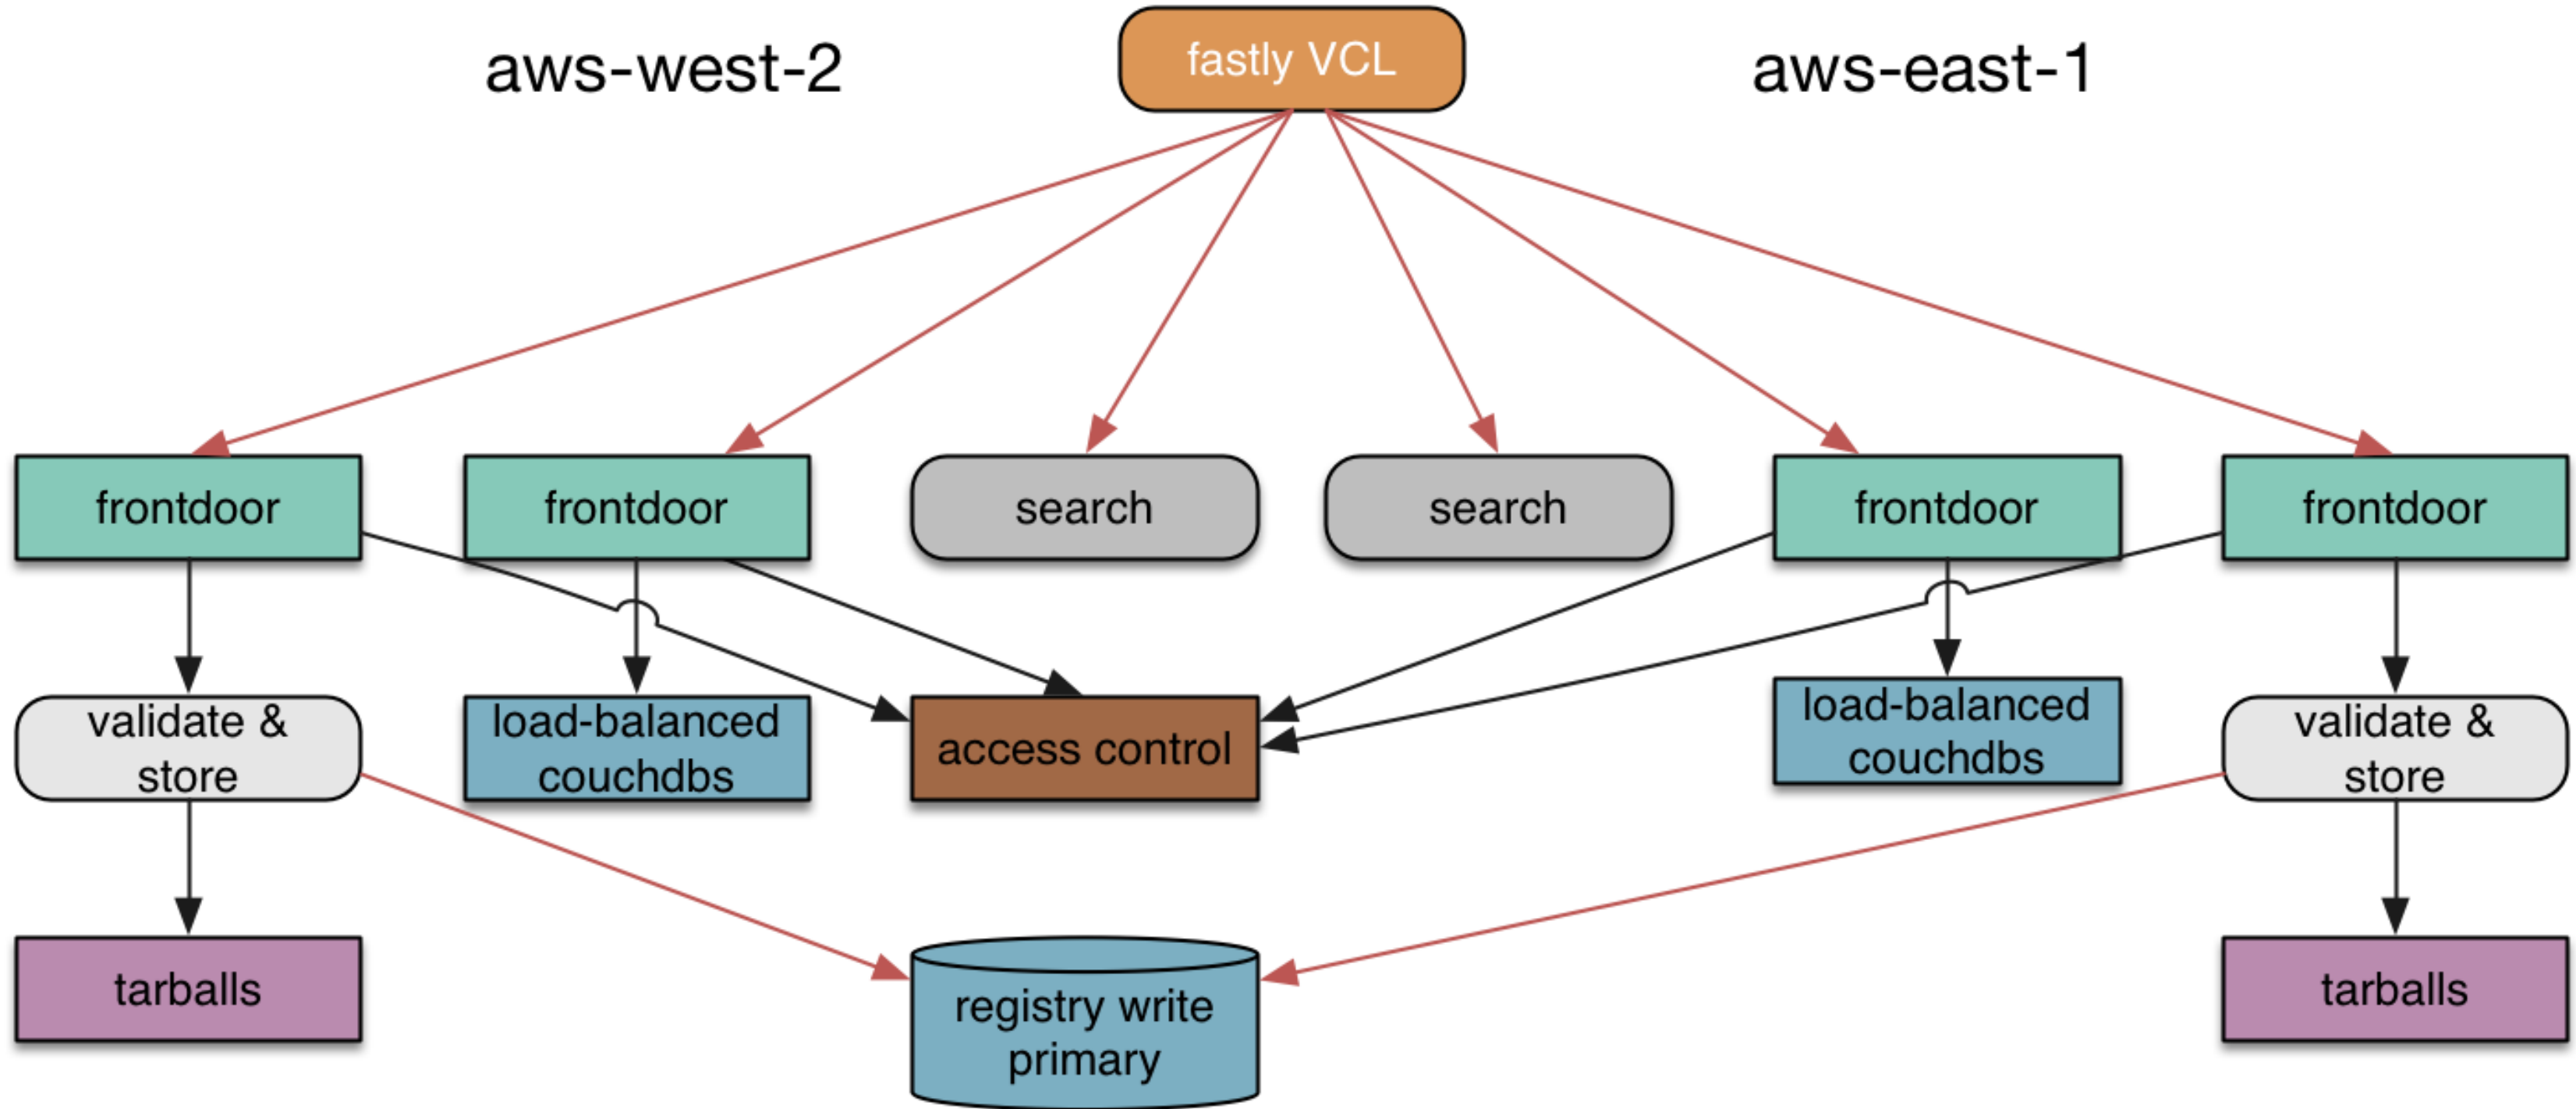
brace yourselves
diagrams incoming

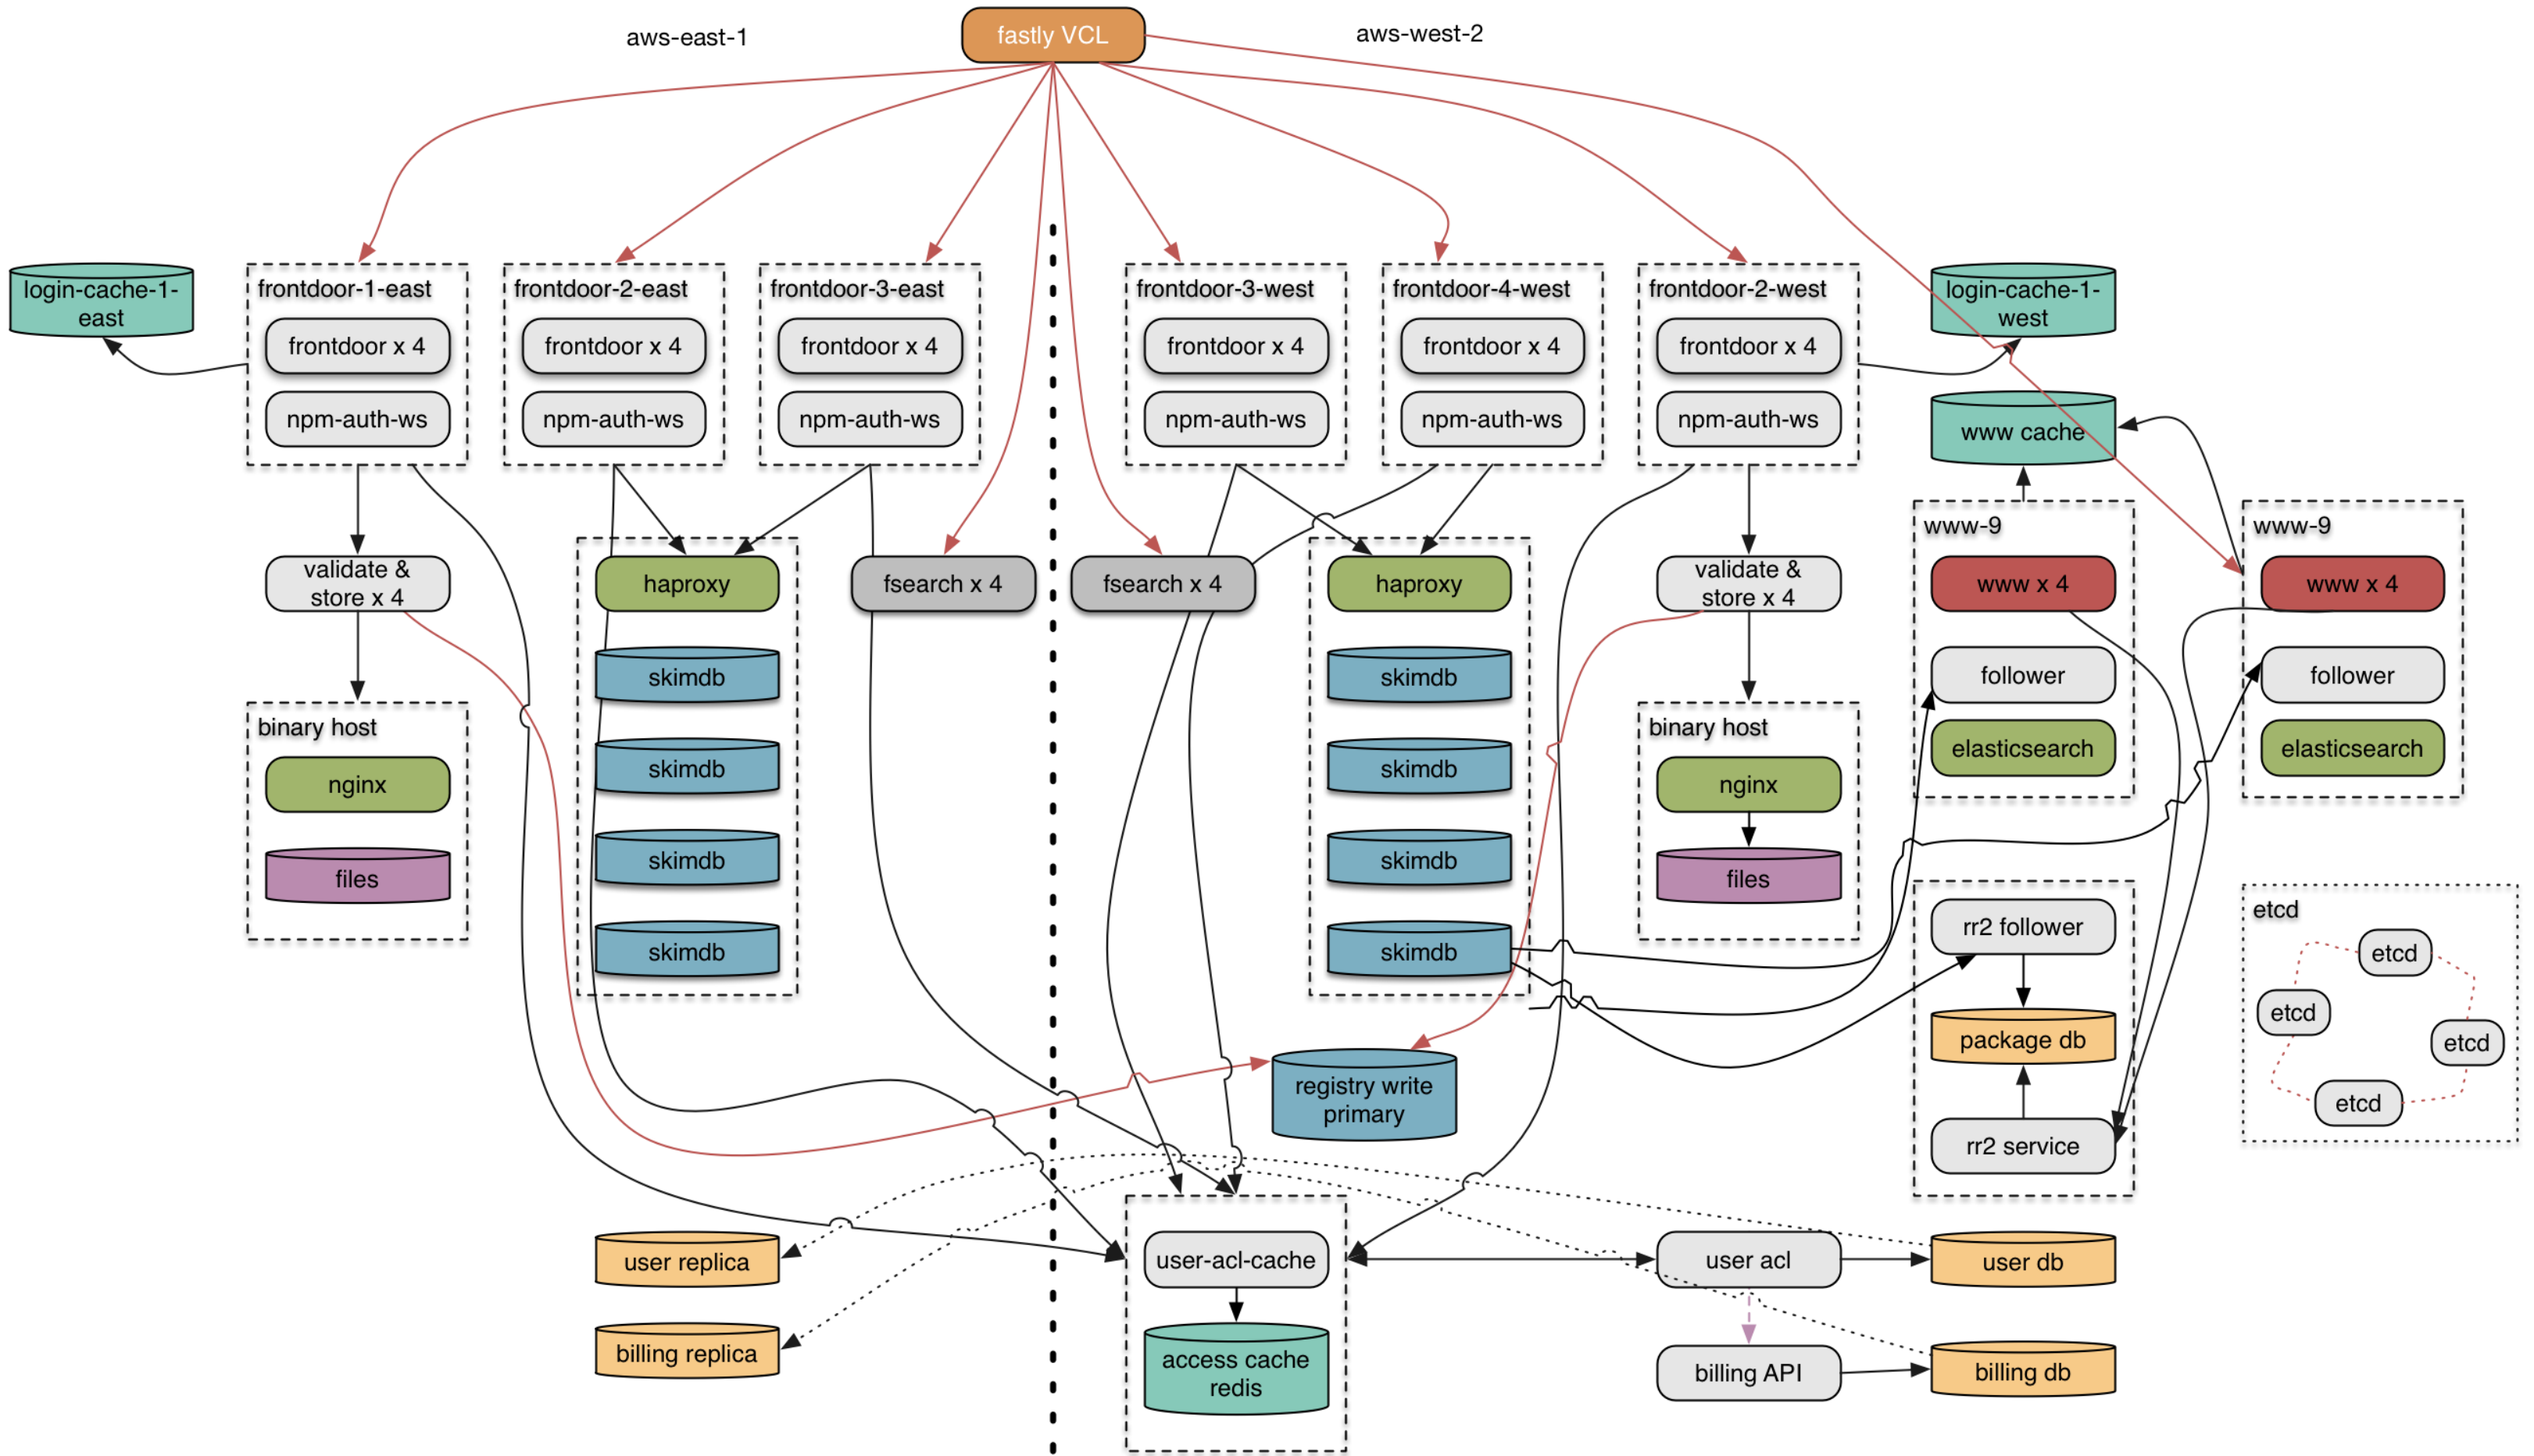


aws-west-2

fastly VCL

aws-east-1





lots of complexity, but

- each piece has a well-defined responsibility**
- each piece can be redundant**
- exceptions: db write primaries**
- each service can be worked on in isolation**

downsides

- yay distributed systems
- pretty sure a message queue is in our future
- some single points of failure: db primaries
- metrics & log handling is poor
- everything is hand-rolled

conservatism won with node

- we're mostly on node 0.10.38**
- memory leaks, some networking trouble with early iojs**
- will try again with iojs 1.8.x**
- or with node now that iojs took over :)**

git deploy

**This was a pain until we wrote a bunch of tools.
Ansible to set it up once. Git to deploy. (Not the
@mafintosh future!)**

```
git push origin +master:deploy-production  
git push origin +master:deploy-staging
```

**Each interested host will report in Slack when it's
done. You've deployed!**

A git-deployable service

- haproxy load-balancing & monitoring
- webhooks server
- github webhooks trigger a bash script
- any server can have many apps git-deployed to it
- generally 1 process per core

open sourced parts

- **jthooks**: set up github web hooks from the command line
- **jthooper**: a server that listens for webhook pushes from github & runs scripts in response
- **rderby**: rolling restarts for servers behind haproxy
- **renv**: recursively manages json blobs with etcd.
- **ndm**: generate upstart/whatever scripts from a service.json config

metrics

All open-source. InfluxDB → Grafana for dashboards.

- numbat-emitter - client to emit metrics from any node service**
- numbat-collector - service to collect & redirect to many outputs**

150,000 modules

~400GB tarballs

68 million dls/day peak

5800 req/sec peak

future work

- organizations for private modules! already in progress**
- make web site search a lot better**
- make the relational package data available via public api**
- more public replication points (all public packages, including scoped)**

npm loves you

```
npm install -g npm@latest
```